TASK 4: Data Warehouse Optimizations

Logical indexed view

Personal notes:

In data warehouses, you can use materialized views **to precompute and store aggregated data such as the sum of sales**.

The SELECT list in the materialized view definition needs to meet at least one of these two criteria: The SELECT list contains an aggregate function.
GROUP BY is used in the Materialized view definition and all columns in GROUP BY are included in the SELECT list.
Aggregate functions are required in the SELECT list of the materialized view definition.
Supported aggregations include MAX, MIN, AVG, COUNT, COUNT_BIG, SUM, VAR, STDEV.

—

For the logical indexed view I chose the query for my **third analytical question**, which is “Which delay type has the longest average delay ?“.

Reason for choosing this query is that it **contains an aggregated calculation** (AVG) and it would be beneficial to not have to query every single delay row when requesting this information, since there are 632 thousand delays (currently).

```
SELECT ddt.delay_code, ddt.delay_description,
AVG(fd.EFFECTIVE_DELAY_TIME_MV) AS average_delay,
COUNT(ddt.delay_code) as delay_count

FROM [datawarehouse].[dbo].[fact_delay] fd
INNER JOIN [datawarehouse].[dbo].[dim_delay_type] ddt
ON fd.DIM_DELAY_TYPE_FK = ddt.delay_type_sk

GROUP BY ddt.delay_code, ddt.delay_description
```

The script:

https://docs.google.com/document/d/1LUMJH2-f0gq6DbV1lXE7WqB30qmwrX-q5adBfuMQeWA/edit?usp=sharing

Or in the Task 4 folder.

Execution plans side by side

Top - With logical indexed view.
Bottom - Without logical indexed view.

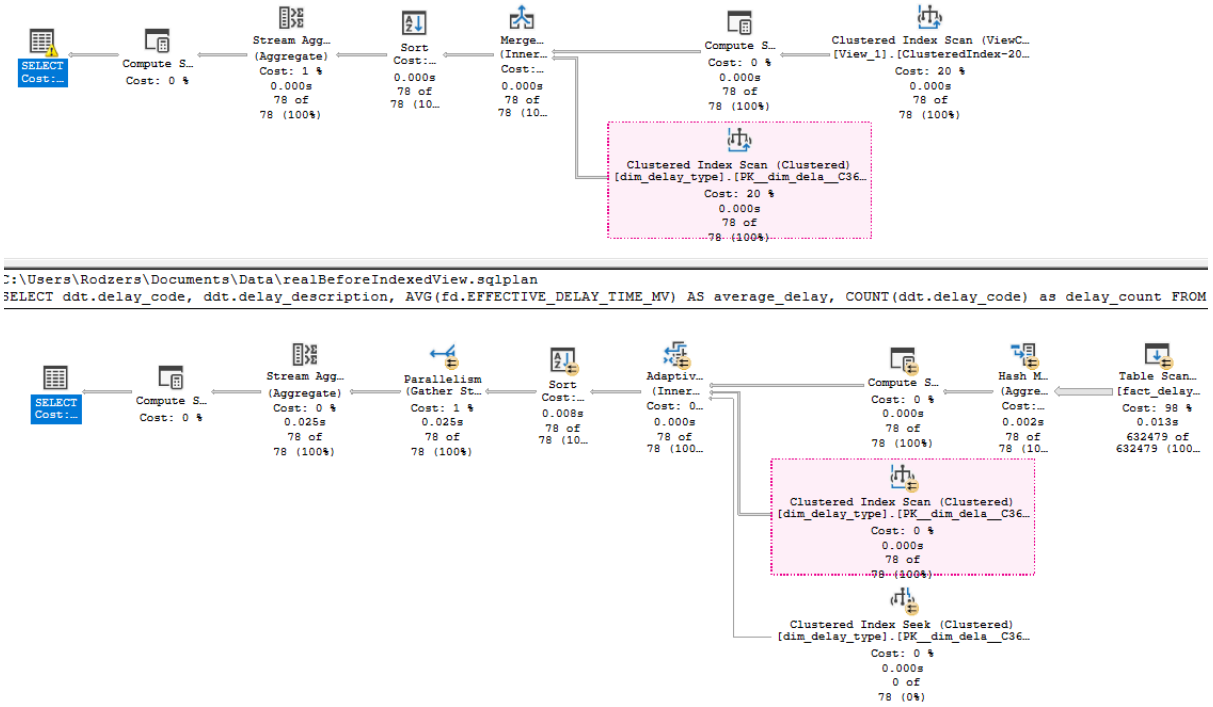

Conclusion

With the help of the logical indexed view, I was able to eliminate the most resource intensive operation, which was the fact_delay table scan.

It’s relative cost to the whole query was 98% and it was due to the fact that it had to read the entire table which at the time of querying was 632’479 rows.

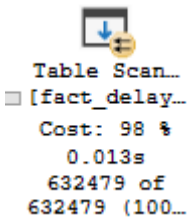

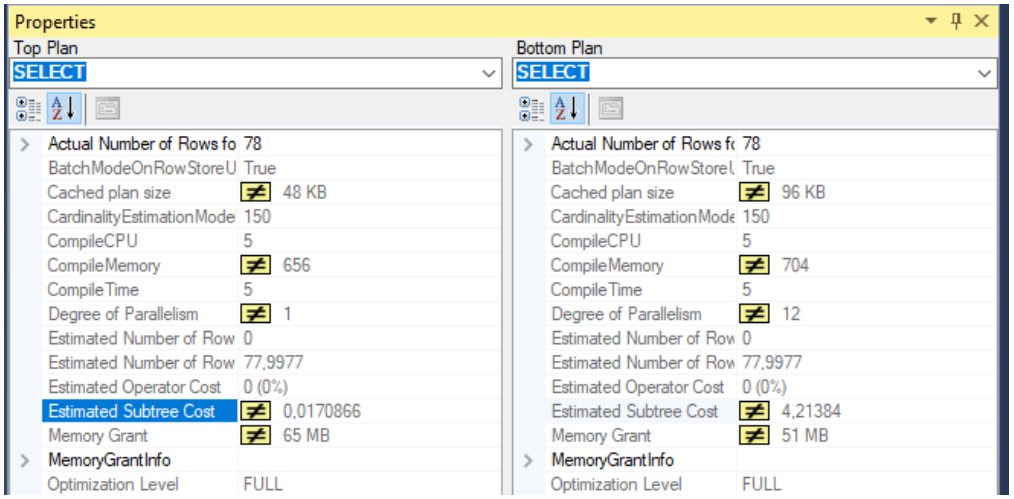

Thanks to this optimization, we can see that the overall query cost has dropped from 4.21384 to 0.0170866, which in my opinion is a **massive** improvement.

And the host resources were reduced as well with the Degree of Parallelism dropping from 12 all the way to 1.

And that is because the table_scan operation was replaced with my newly created Clustered Index.

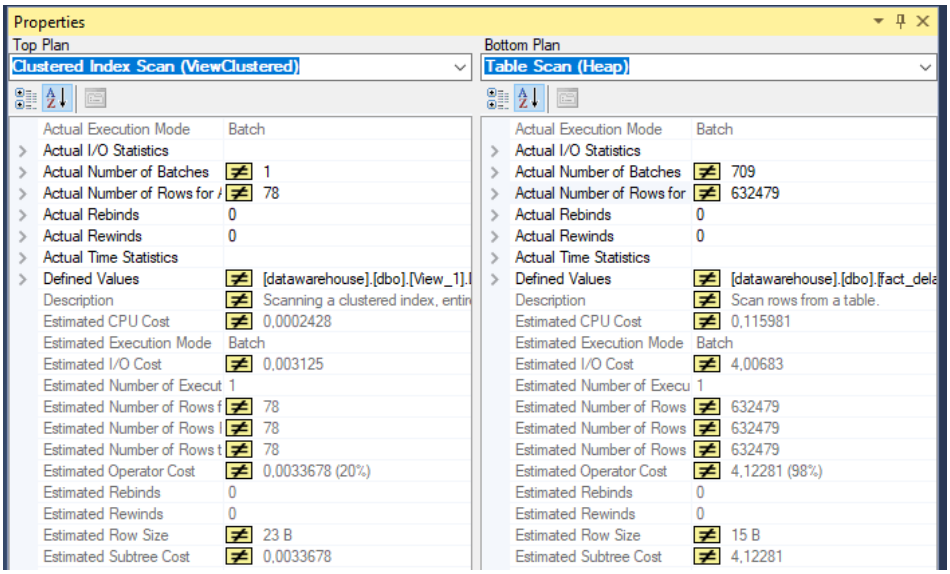

Some notable observations:

- Number of rows that it has to scan is now 78 instead of 632’479.
- Estimated CPU Cost dropped from 0.115981 to 0.0002428.
- Estimated I/O Cost dropped from 4.0683 to 0.003125.
- And the Estimated Operator Cost for acquiring the necessary data is now 0.0033678 instead of 4.12281

Would I implement this logical indexed view ?

That depends on how often I will use this query and whether this information is vital to my ‘business’ / ‘company’.

If it’s only being used once a month (or even less frequently) for preparing presentations / summaries within the company, then there’s no reason in my opinion to waste additional disk space and resources.

Otherwise if the query serves some valid purpose, then why not, the benefits would outweigh the disadvantages.

Column storage

Personal notes:

Column storage is optimized for fast retrieval of columns of data, which is useful for analytical applications, due to the fact that it drastically reduces the overall disk I/O requirements and reduces the amount of data that you need to load from disk.

It’s often used in data warehouses where it’s necessary to send large amounts of data from multiple sources for these previously mentioned BI applications.

It was designed to accelerate data warehousing queries, which require scanning, aggregation and filtering of large amounts of data, or joining multiple tables like a star schema.

Requirements:

- Limited amount of columns.

Therefore the best place to implement this columnstore storage is in the **fact_delay** table of my datawarehouse.

The script:

https://docs.google.com/document/d/1FDfXJOAlcQTMtr40EqH3UJzzjqGn5uDoDT5-xqlE4WQ/edit?usp=sharing

Or in the Task 4 folder.

Execution plans side by side

Top - With column store index.
Bottom - Without column store index.

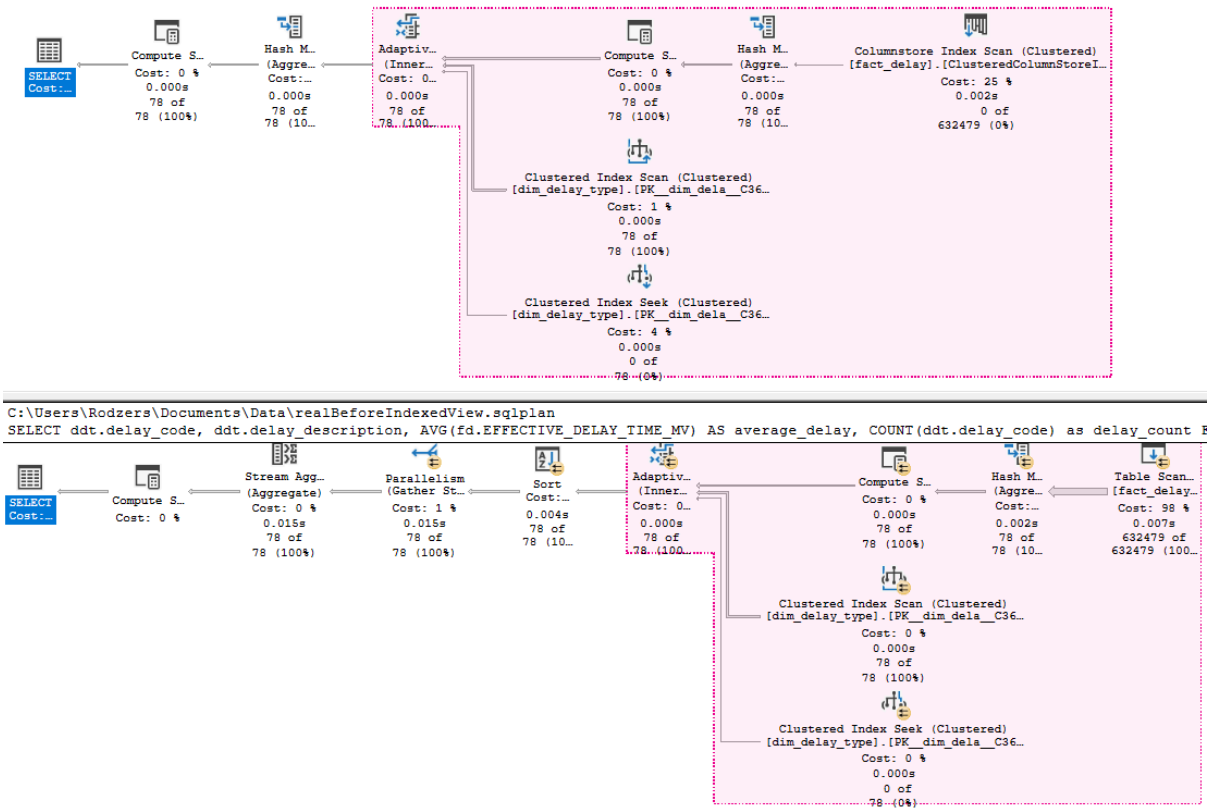

Conclusion

Once again you can see large overall improvements.

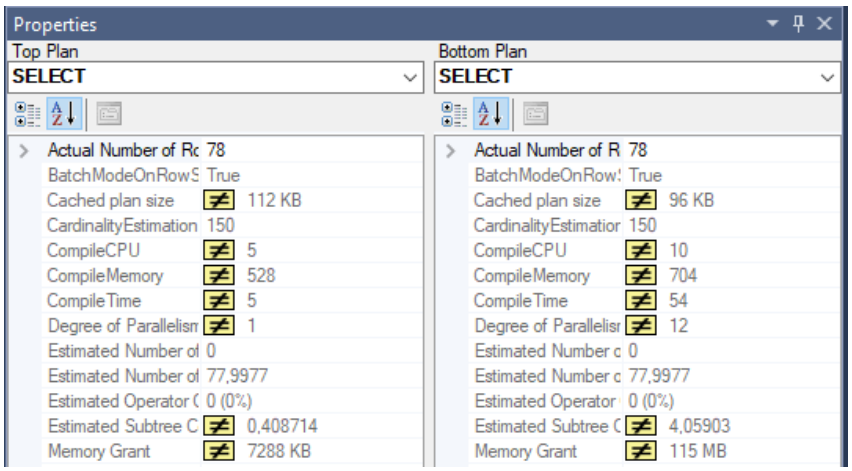

But not as much as for the Logical Indexed View.

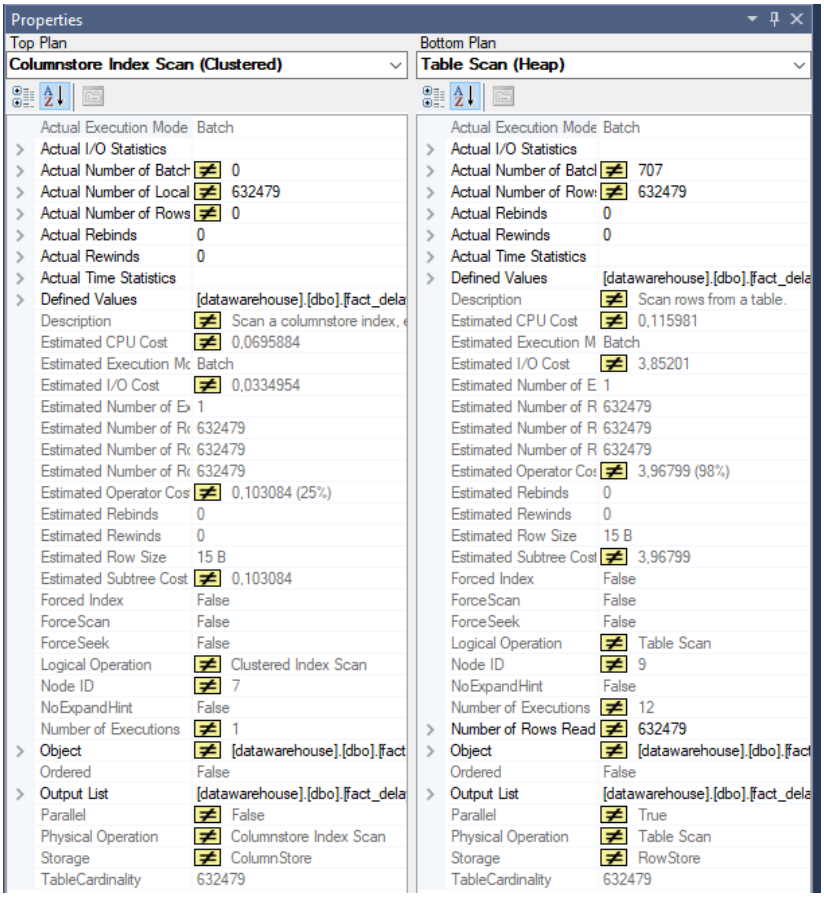

The Estimated Operator Costs dropped from 3.96799 to 0.103084.

And this time you can also see massive improvements in the memory sector, where the query execution only required 7168 KB of memory instead of 51176 KB and used 1024 KB instead of 6464.

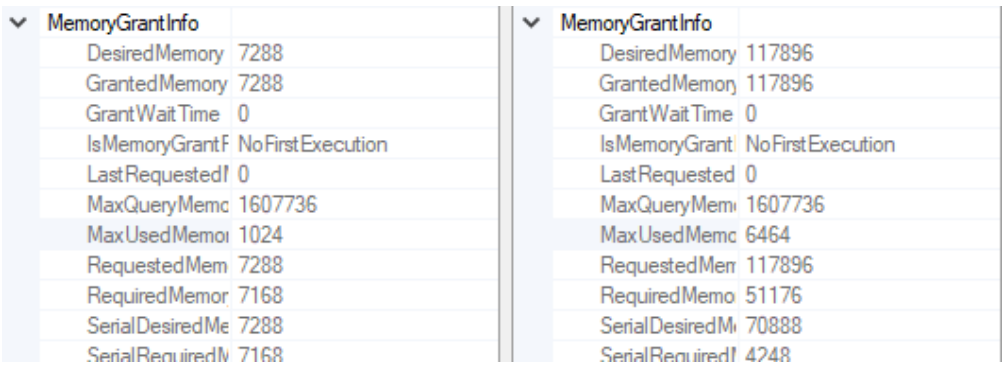

Would I implement the column storage specifically for this query ?

For this specific query, no, but overall for an OLAP system, more specifically my fact_delay table, yes.